


Concurrency

5 - Monitors & Condition Synchronization




Credits for the slides:
Claus Braband
Jeff Magee & Jeff Kramer

Alexandre David
adavid@cs.aau.dk

Repetition - Interference (Ornamental Garden Problem)

People enter an ornamental garden through either of two turnstiles. Management wishes to know how many are in the garden at any time. (Nobody can exit).



Repetition - Running the Applet



After the East and West turnstile threads each have incremented the counter 20 times, the garden people counter is not the sum of the counts displayed.

Repetition - Model Checking (reveals the error)

Ornamental Garden Model reveals the error:

```
|| TESTGARDEN = ( GARDEN || TEST ).
```

- Use *LTSA* to perform an exhaustive search for **ERROR**:

Trace to property violation in TEST:

```
go
east.arrive
east.value.read.0
west.arrive
west.value.read.0
east.value.write.1
west.value.write.1
end
display.value.read.1
wrong
```

LTSA produces
the shortest
path to reach
the **ERROR** state.

Repetition - Interference and Mutual Exclusion

- ◆ Interference (Java):

```
x = x + 1; || x = x + 1;
```


- ◆ Mutual exclusion (Java):

```
synchronized (obj) {  
    x = x + 1;  
}
```

```
synchronized (obj) {  
    x = x + 1;  
}
```

- ◆ Modelling mutual exclusion (FSP):

```
LOCK = (acq -> rel -> LOCK).
```



Monitors & Condition Synchronization

Concepts: monitors:

encapsulated data + access procedures

mutual exclusion + condition synchronization

single access procedure active in the monitor

nested monitors


Models: guarded actions

Practice: private data and synchronized methods (exclusion).

wait(), notify() and notifyAll() for condition synch.

single thread active in the monitor at a time

5.1 Condition Synchronization (Car Park)



A controller is required to ensure:

- cars can only enter when not full
- cars can only leave when not empty (duh!)

Car Park Model (Actions and Processes)



◆ Actions of interest:

- **arrive**
- **depart**

◆ Identify processes:

- Arrivals
 - Departures
 - Control
- $\left. \begin{array}{l} \bullet \text{Arrivals} \\ \bullet \text{Departures} \\ \bullet \text{Control} \end{array} \right\} env$


Car Park Model (Structure Diagram)

◆ Actions of interest:

- arrive
- depart

◆ Identify processes:

- Arrivals
 - Departures
 - Control
- } *env*



Concurrency: monitors & condition synchronization

©Magee/Kramer

Car Park Model (FSP)

```
ARRIVALS = (arrive -> ARRIVALS).  
  
DEPARTURES = (depart -> DEPARTURES).  
  
CONTROL(N=4) = SPACES[N],  
  
SPACES[i:0..N] = (when(i>0) arrive -> SPACES[i-1]  
                  | when(i<N) depart -> SPACES[i+1]).  
  
|| CARPARK = (ARRIVALS || DEPARTURES || CONTROL(4)).
```

Guarded actions are used to control `arrive` and `depart`
LTS?


Car Park Program

- ◆ Model
 - ◆ - all entities are processes interacting via *shared actions*
- ◆ Program - need to identify threads and monitors:
 - ◆ thread - **active** entity which initiates (output) actions
 - ◆ monitor - **passive** entity which responds to (input) actions.

For the carpark?

- | | | | |
|---------------|----------------|----|---------|
| • Arrivals: | active | => | thread |
| • Departures: | active | => | thread |
| • Control: | passive | => | monitor |

Car Park Program (Interesting part of Class Diagram)



Car Park Program - Applet::start()

The Applet's start() method creates:

- **CarParkControl** monitor (with condition synchr.)
- **Arrival** thread
- **Departures** thread

```
public void start() {  
    CarParkControl c = new DisplayCarPark(disp, PLACES);  
    arrivals.start(new Arrivals(c));  
    departures.start(new Departures(c));  
}
```

The **CarParkControl** is *shared* by **Arrival** and **Departures** threads

Car Park Program - Arrivals and Departures threads

```
class Arrivals implements Runnable {  
    CarParkControl carpark;  
  
    Arrivals(CarParkControl c) { carpark = c; }  
  
    public void run() {  
        try {  
            while(true) {  
                ThreadPanel.rotate(330);  
                carpark.arrive();  
                ThreadPanel.rotate(30);  
            }  
        } catch (InterruptedException _) {}  
    }  
}
```

Similarly,
Departures calls:

carpark.depart()

How do we implement the control of **CarParkControl**?

Car Park Program - CarParkControl Monitor

```
class CarParkControl {  
    protected int spaces, capacity;  
  
    CarParkControl(int n) {  
        capacity = spaces = n;  
    }  
  
    synchronized void arrive() {  
        ... --spaces; ...  
    }  
  
    synchronized void depart() {  
        ... ++spaces; ...  
    }  
}
```

Encapsulation
~ **protected**

Mutual exclusion
~ **synchronized**

Condition synchronization?

Block if full?
($spaces == 0$)

Block if empty?
($spaces == N$)

Condition Synchronization in Java

Java provides a **thread wait queue per object** (*not per class*).

Object has methods:

```
public final void wait() throws InterruptedException;
```

Waits to be *notified* (i.e. another thread invokes **notify**).
Releases the synchronization lock associated with the obj.

When notified, the thread must reacquire the synchr. lock.

```
public final void notify();
```

```
public final void notifyAll();
```


Wakes up (notifies) a thread waiting on the object's queue.

Condition Synchronization in Java (enter/exit)

A thread:

- *Enters* a monitor when a thread acquires the lock associated with the monitor;
- *Exits* a monitor when it releases the lock.

Wait() causes the thread to **exit** the monitor, permitting other threads to **enter** the monitor



Condition Synchronization in FSP and Java

FSP: when(*cond*) *action->NEWSTATE*

```
synchronized void act() throws InterruptedException {  
    while (!cond) wait();  
    // modify monitor data  
    notifyAll();  
}
```

The `while` loop is necessary to re-test the condition *cond* to ensure that *cond* is indeed satisfied when it re-enters the monitor.

`notifyAll()` is necessary to awaken other thread(s) that may be waiting to enter the monitor now that the monitor data has been changed.

CarParkControl - Condition Synchronization

```
class CarParkControl {  
    protected int spaces, capacity;  
  
    synchronized void arrive() throws Int'Exc' {  
        while (spaces==0) wait();  
        --spaces;  
        notify();  
    }  
  
    synchronized void depart() throws Int'Exc' {  
        while (spaces==capacity) wait();  
        ++spaces;  
        notify();  
    }  
}
```

Why is it sensible to use `notify()` here rather than `notifyAll()`?

Models to Monitors - Guidelines


- **Active** entities (that initiate actions) are implemented as **threads**.
- **Passive** entities (that respond to actions) are implemented as **monitors**.

Each guarded action in the model of a monitor is implemented as a **synchronized** method which uses a **while** loop and **wait()** to implement the guard.

The **while** loop condition is the negation of the model guard condition.

Changes in the state of the monitor are signaled to waiting threads using **notifyAll()** (or **notify()**).

5.2 Semaphores



Semaphores are widely used for dealing with inter-process synchronization in operating systems.

Semaphore s : integer var that can take only non-neg. values.

$s.\text{down}()$: when $s > 0$ do $\text{decrement}(s)$; Aka. "P" ~ Passern

$s.\text{up}()$: $\text{increment}(s)$; Aka. "V" ~ Vrijgeven

Usually implemented as *blocking wait*:

$s.\text{down}()$: if ($s > 0$) then $\text{decrement}(s)$;
else block execution of calling process

$s.\text{up}()$: if (processes blocked on s) then awake one of them
else $\text{increment}(s)$;

Modelling Semaphores


To ensure analyzability, we only model semaphores that take a finite range of values. If this range is exceeded then we regard this as an ERROR. N is the initial value.

```
const Max = 3
range Int = 0..Max

SEMAPHORE(N=0) = SEMA[N],
SEMA[v:Int]     = (up->SEMA[v+1]
                     | when(v>0) down->SEMA[v-1]),
SEMA[Max+1]     = ERROR.
```

LTS?

Modelling Semaphores



Action down is only accepted when value (v) of the semaphore is greater than 0.

Action up is not guarded.

Trace to a violation:

up → up → up → up

Semaphore Demo - Model

Three processes $p[1..3]$ use a shared semaphore `mutex` to ensure mutually exclusive access (action `critical`) to some resource.

```
LOOP = (mutex.down->critical->mutex.up->LOOP).  
|| SEMADEMO = (p[1..3]:LOOP  
||   {p[1..3]}::mutex:SEMAPHORE(1)).
```

For mutual exclusion, the semaphore initial value is 1. *Why?*


Is the ERROR state reachable for SEMADEMO?

Is a binary semaphore sufficient (i.e. Max=1) ?

LTS?



Semaphore Demo - Model




Semaphores in Java

Semaphores: passive objects => implemented as monitors.

```
public class Semaphore {  
    private int value;  
  
    public Semaphore (int n) { value = n; }  
  
    synchronized public void up() {  
        ++value;  
        notify();  
    }  
  
    synchronized public void down() throws Int'Exc' {  
        while (value == 0) wait();  
        --value;  
    }  
}
```

*In practice,
semaphore is a low-level mechanism often used in
implementing higher-level monitor constructs.*

SEMADEMO display



current
semaphore
value

thread 1 is
executing
critical
actions.

thread 2 is
blocked
waiting.

thread 3 is
executing
non-critical
actions.

SEMADEMO

What if we adjust the time that each thread spends in its critical section ?

◆ large resource requirement - *more conflict?*

(eg. more than 67% of a rotation)?

◆ small resource requirement - *no conflict?*

(eg. less than 33% of a rotation)?

Hence the time a thread spends in its critical section should be kept as short as possible.

SEMADEMO Program - MutexLoop

```
class MutexLoop implements Runnable {
    Semaphore mutex;

    MutexLoop (Semaphore sema) {mutex=sema; }

    public void run() {
        try {
            while(true)  {
                while(!ThreadPanel.rotate());
                mutex.down();           // acquire
                while(ThreadPanel.rotate()); // critical
                mutex.up();             // release
            }
        } catch(InterruptedException _) {}
    }
}
```


Threads and semaphore are created by the applet start() method.

ThreadPanel.rotate() returns false while executing non-critical actions (dark color) and true otherwise.

5.3 Bounded Buffer


A bounded buffer consists of a fixed number of slots.

Items are put into the buffer by a *producer* process and removed by a *consumer* process:




≈ Car Park Example!

Bounded Buffer - a Data-Independent Model



The *behaviour* of BOUNDED BUFFER is independent of the actual data values, and so can be modelled in a data-independent manner.

LTS?



Bounded Buffer - a Data-Independent Model

```
PRODUCER = (put->PRODUCER).  
  
CONSUMER = (get->CONSUMER).  
  
BUFFER(N=5) = COUNT[ 0 ],  
COUNT[i:0..N] = (when (i<N) put->COUNT[i+1]  
                 |when (i>0) get->COUNT[i-1]).  
  
|| BOUNDED BUFFER =  
( PRODUCER || BUFFER( 5 ) || CONSUMER ).
```

Bounded Buffer Program - Buffer Monitor

```
public interface Buffer {  
    public void put(Object o) throws InterruptedException;  
    public Object get()           throws InterruptedException;  
}
```

```
class BufferImpl implements Buffer {  
    protected Object[] buf;  
    protected int in, out, count, size;  
    ...  
    synchronized void put(Object o) throws InterruptedException {  
        while (count==size) wait();  
        buf[in] = o;  
        count++;  
        in = (in+1) % size;  
        notifyAll();  
    }  
}
```

We separate the interface to permit an alternative implementation later.

Similarly for get()

```
synchronized Object get() throws Int'Exc' {  
1.     while (count==0) wait();  
2.     Object obj = buf[out];  
3.     buf[out] = null;  
4.     count--;  
5.     out = (out+1) % size;  
6.     notifyAll();  
7.     return obj;  
}
```

- What happens if we move `notifyAll()` up earlier (e.g. after line 1)?
- What is the point of line 3?

Bounded Buffer Program - Producer Process

```
class Producer implements Runnable {  
    Buffer buf;  
    String alpha = "abcdefghijklmnopqrstuvwxyz";  
  
    Producer(Buffer b) { buf = b; }  
  
    public void run() {  
        try {  
            int i = 0;  
            while(true) {  
                ThreadPanel.rotate(12);  
                buf.put(new Character(alpha.charAt(i)));  
                i=(i+1) % alpha.length();  
                ThreadPanel.rotate(348);  
            }  
        } catch (InterruptedException _) {}  
    }  
}
```

Similarly Consumer
which calls buf.get()

5.4 Nested Monitors

Suppose that, instead of using the *count* variable and condition synchronization, we instead use 2 semaphores *full* and *empty* to reflect the state of the buffer:

```
class SemaBuffer implements Buffer {  
    protected Object buf[];  
    protected int in, out, count, size;  
  
    Semaphore full;    //counts number of items  
    Semaphore empty;  //counts number of spaces  
  
    SemaBuffer(int s) {  
        size = s; in = out = count = 0;  
        buf = new Object[size];  
        full = new Semaphore(0);  
        empty = new Semaphore(size);  
    }  
}
```

Nested Monitors - Bounded Buffer Program

```
synchronized public void put(Object o) throws Int'Exc' {
    empty.down();
    buf[in] = o;
    count++;
    in = (in+1) % size;
    full.up();
}
```

empty is decremented during a **put**,
which is blocked if *empty* is zero.

```
synchronized public Object get() throws Int'Exc' {
    full.down();
    Object o = buf[out];
    buf[out] = null;
    count--;
    out = (out+1) % size;
    empty.up();
    return o;
}
```

full is decremented by a **get**,
which is blocked if *full* is zero.

Does this behave as desired?

Nested Monitors - Bounded Buffer Model

```
PRODUCER = (put -> PRODUCER).  
  
CONSUMER = (get -> CONSUMER).  
  
SEMAPHORE(N=0) = SEMA[N],  
SEMA[v:Int]      = (up->SEMA[v+1]  
                     | when(v>0) down->SEMA[v-1]).  
  
BUFFER = (put -> empty.down -> full.up -> BUFFER  
         | get -> full.down -> empty.up -> BUFFER).  
  
|| BOUNDEDBUFFER =  
  ( PRODUCER || BUFFER || CONSUMER  
    || empty:SEMAPHORE(5)  
    || full:SEMAPHORE(0) ).
```

Does this behave as desired?

Nested Monitors - Bounded Buffer Model

*LTS*A analysis predicts a possible DEADLOCK:

Composing
potential DEADLOCK

States Composed: 28 Transitions: 32 in 60ms

Trace to DEADLOCK:

get

The Consumer tries to get a character, but the buffer is empty. It blocks and releases the lock on the semaphore full. The Producer tries to put a character into the buffer, but also blocks. *Why?*

Nested Monitors - Bounded Buffer Model

*LTS*A analysis predicts a possible DEADLOCK:

Composing

potential DEADLOCK

States Composed: 28 Transitions: 32 in 60ms

Trace to DEADLOCK:

get

- 1) Consumer calls SemaBuffer.get(), acquiring a lock on the buffer
`synchronized public Object get()`
- 2) Semaphore.down() acquires another lock on the Semaphore
`synchronized public void down()`
- 3) Semaphore.down() releases *only its own* lock using wait()
- 4) Producer calls SemaBuffer.put(), blocking on the buffer
`synchronized public void put(Object)`

This situation is known as the *nested monitor problem*.

Nested Monitors - Revised Bounded Buffer Program

The only way to avoid it in Java is by *careful design*. In this example, the deadlock can be removed by ensuring that the monitor lock for the buffer is not acquired until *after* semaphores are decremented.

```
public void put(Object o) throws Int'Exc' {
    empty.down();
    synchronized (this) {
        buf[in] = o;
        count++;
        in = (in+1) % size;
    }
    full.up();
}
```

Nested Monitors - Revised Bounded Buffer Model

```
BUFFER = (put -> BUFFER  
          |get -> BUFFER).
```

```
PRODUCER = (empty.down -> put -> full.up -> PRODUCER).  
CONSUMER = (full.down -> get -> empty.up -> CONSUMER).
```

The semaphore actions have been moved to the producer and consumer. This is exactly as in the implementation where the semaphore actions are *outside* the monitor .

Does this behave as desired?

Minimized LTS?

5.5 Monitor invariants

An invariant for a monitor is an assertion concerning the variables it encapsulates. This assertion must hold whenever there is no thread executing inside the monitor i.e. on thread entry to and exit from a monitor .

INV(CarParkControl): $0 \leq spaces \leq N$

INV(Semaphore): $0 \leq value$

INV(Buffer): $0 \leq count \leq size$

and $0 \leq in < size$

and $0 \leq out < size$

and $in = (out + count) \% size$

Like normal invariants, but must also hold when lock is released (wait)!

Summary

Concepts: monitors:

encapsulated data + access procedures

mutual exclusion + condition synchronization

single access procedure active in the monitor

nested monitors

Models: guarded actions

Practice: private data and synchronized methods (exclusion).

wait(), notify() and notifyAll() for condition synch.

single thread active in the monitor at a time