



Processes in UNIX

Emmanuel Fleury
B1-201
fleury@cs.aau.dk



WIKIPEDIA
The Free Encyclopedia





Overview

- Motivations
 - Problems in Concurrency
 - A Process Life
 - Interlude
 - Process Management
 - Process Scheduling
-

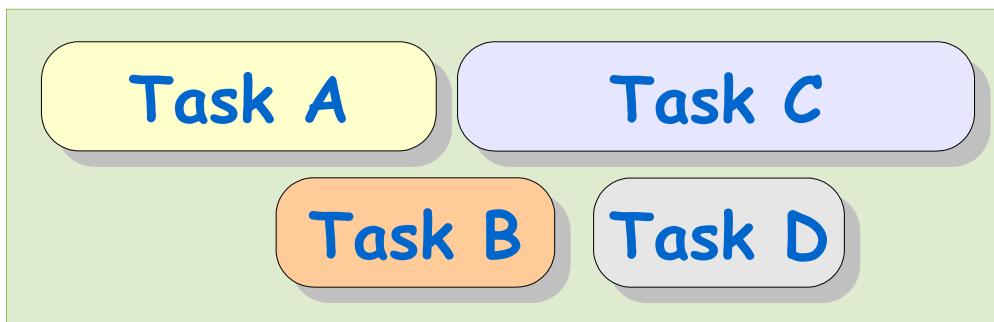
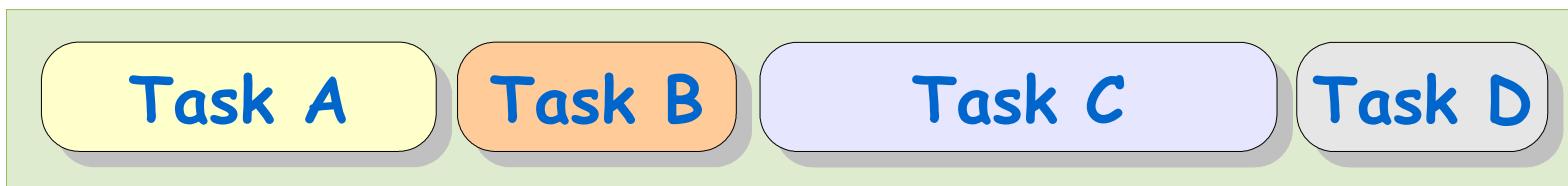


Motivations



Basic Idea

Do several things at once !



Efficiency, Responsiveness, Scalability, ...



Concurrent People...

A natural thing ?



... and Concurrent Machines...

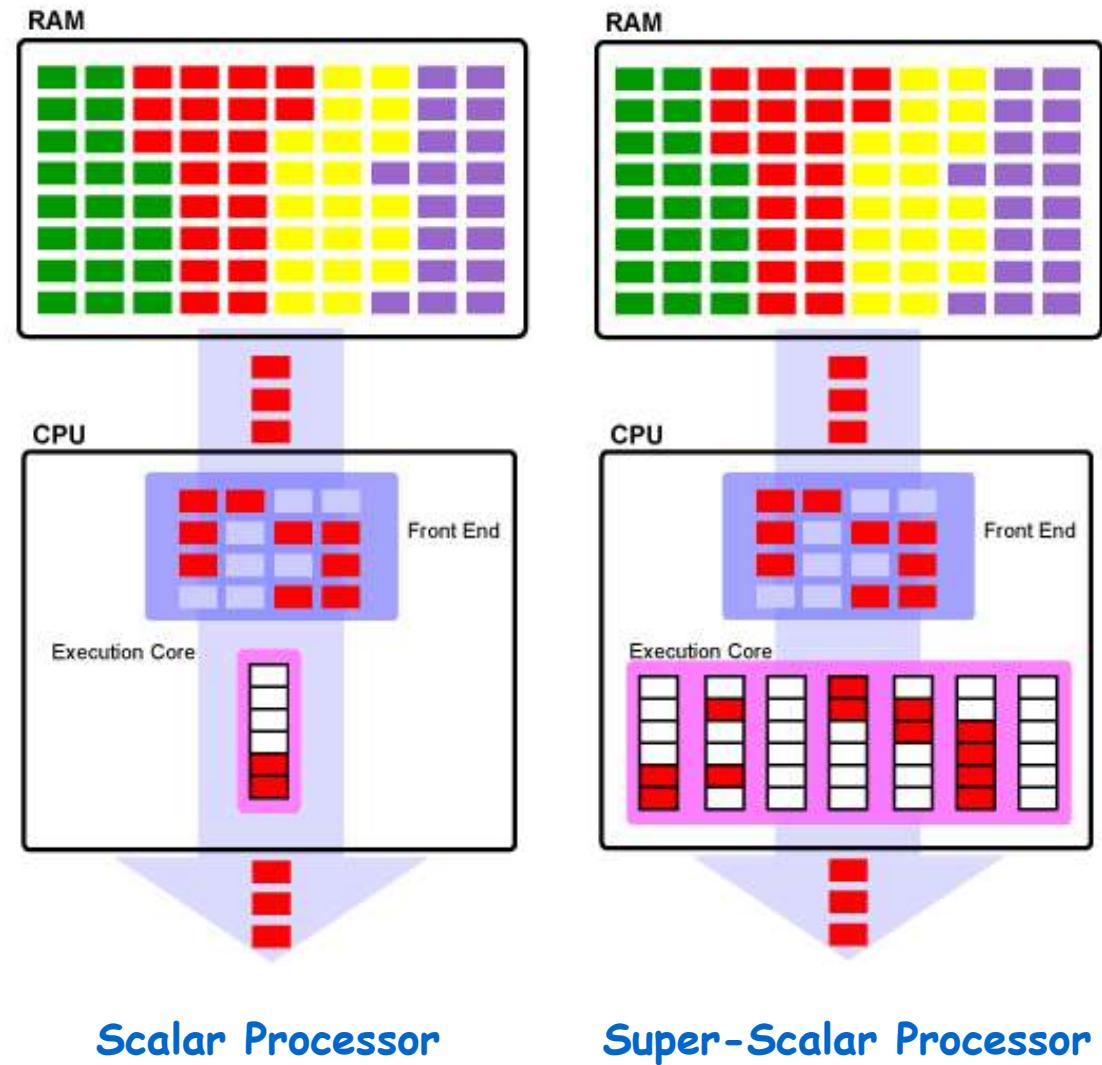
More and more parallelism !!!

- [1970s] Symmetric Multi-Processing (SMP)
- [1980s] Super-scalar Processors
- [1990s] Hyper-threading Processors
- [2000s] Cell Processors
- ... and other will follow



Super-scalar Processors

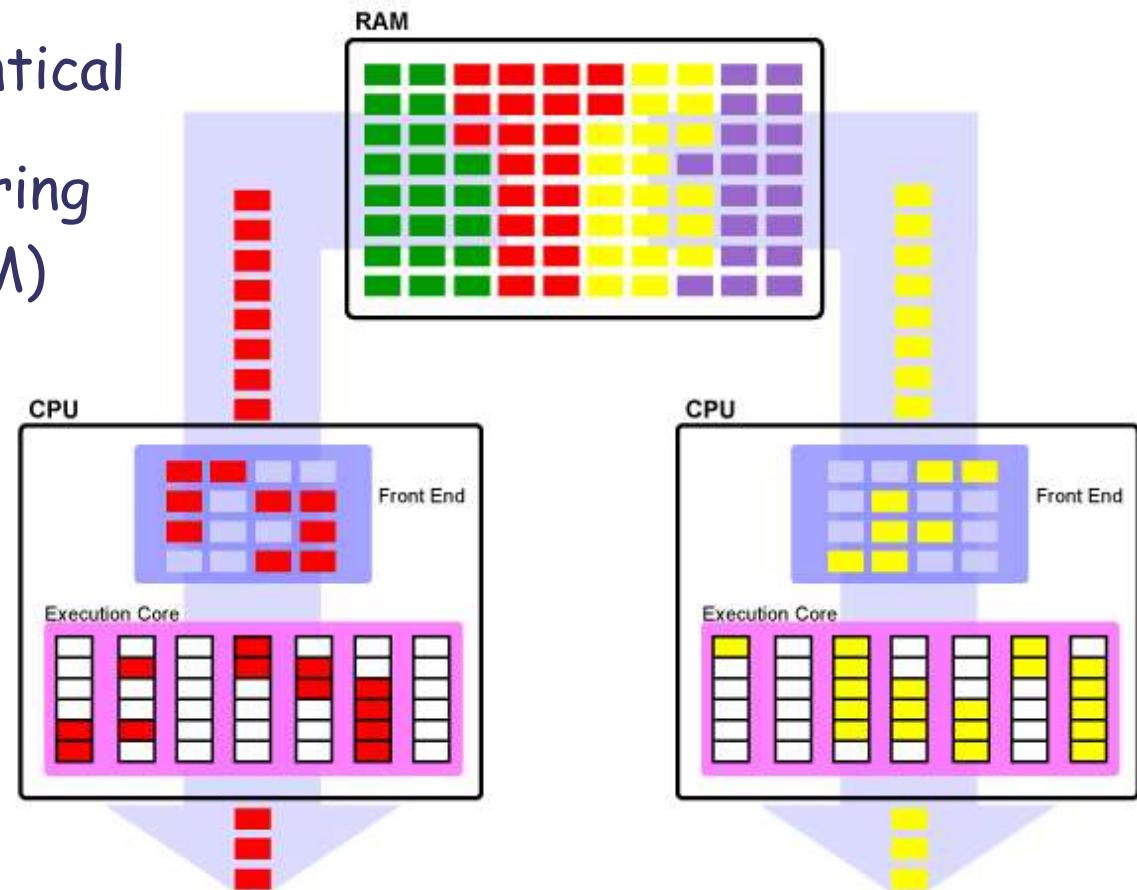
- Implement Parallelism on a single chip
- Dispatch the tasks among several processing units
- Efficiency strongly depends on the dispatcher
- Difficult to compute the WCET (Worst Case Execution Time)





Symmetric Multi-Processing

- Multi-processors machines
- All processors are identical
- All processors are sharing the same memory (RAM)
- Multi-processing is for real !!
- Now a days a lot of desktops use this technology

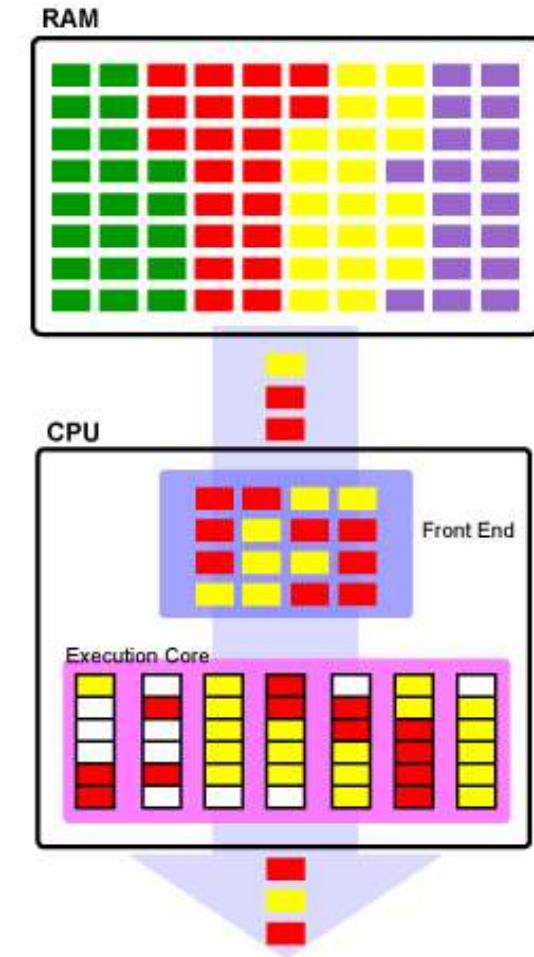


Symmetric Multi-processing



Hyper-Threading

- Emulate two processors in one through the dispatcher
- Improve efficiency and reactivity to multi-processed and/or multi-threaded programs
- Performance improvements of 15-30%
- Tend to be more and more common

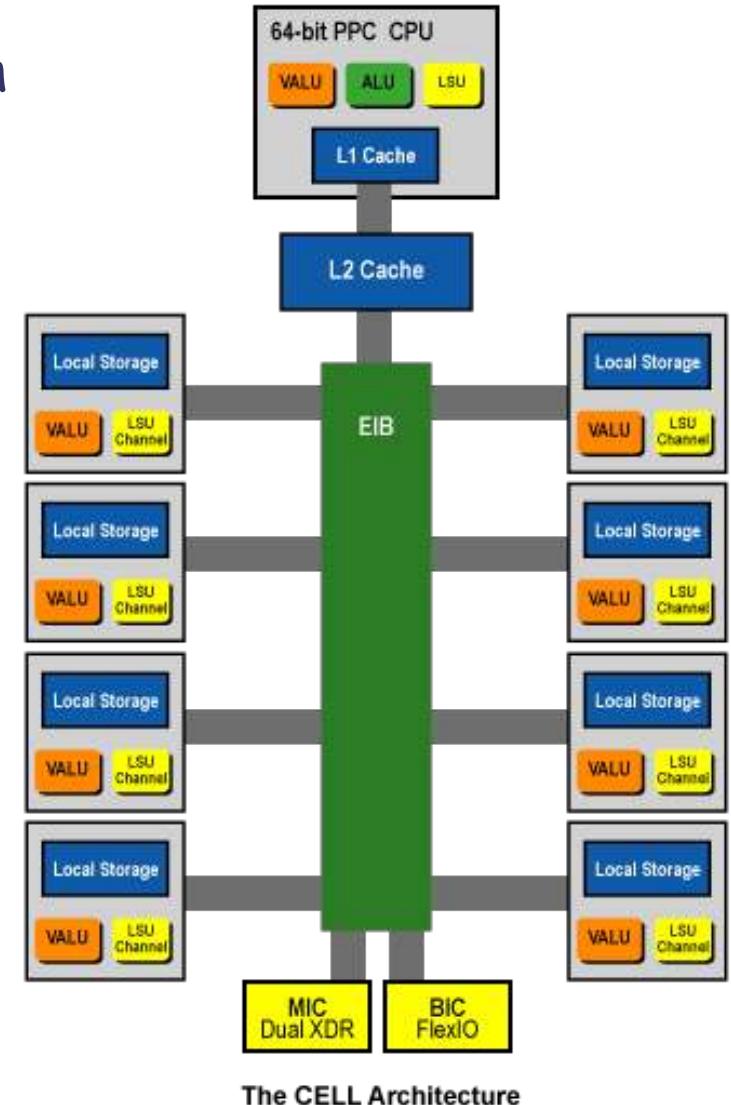


Hyper-Threading



Cell Processors

- Developed by IBM in cooperation with Toshiba and Sony (Playstation 3).
- Highly Parallel Architecture
- Designed for streaming (audio, video)
- Main Components:
 - 1 Processing Element (PE)
 - 1 Element Interconnection Bus (EIB)
 - 8 Synergistic Processing Units (SPU)



The CELL Architecture



... in a Concurrent World !

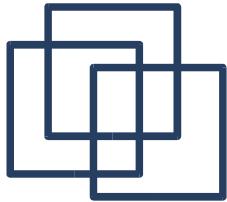
- Applications can spread over networks (Internet)
- Scalability go often through the use of others machines
- Multi-processed and/or Multi-threaded applications can better be adapted for networks



Conclusion ?

Think concurrent !

- The users are more and more familiar with multi-tasking
- The hardware tend to handle more and more efficiently concurrent programs
- The networks are growing in bandwidth and power
- The operating systems as well (Schedulers)



Common Problems in Concurrent Programming



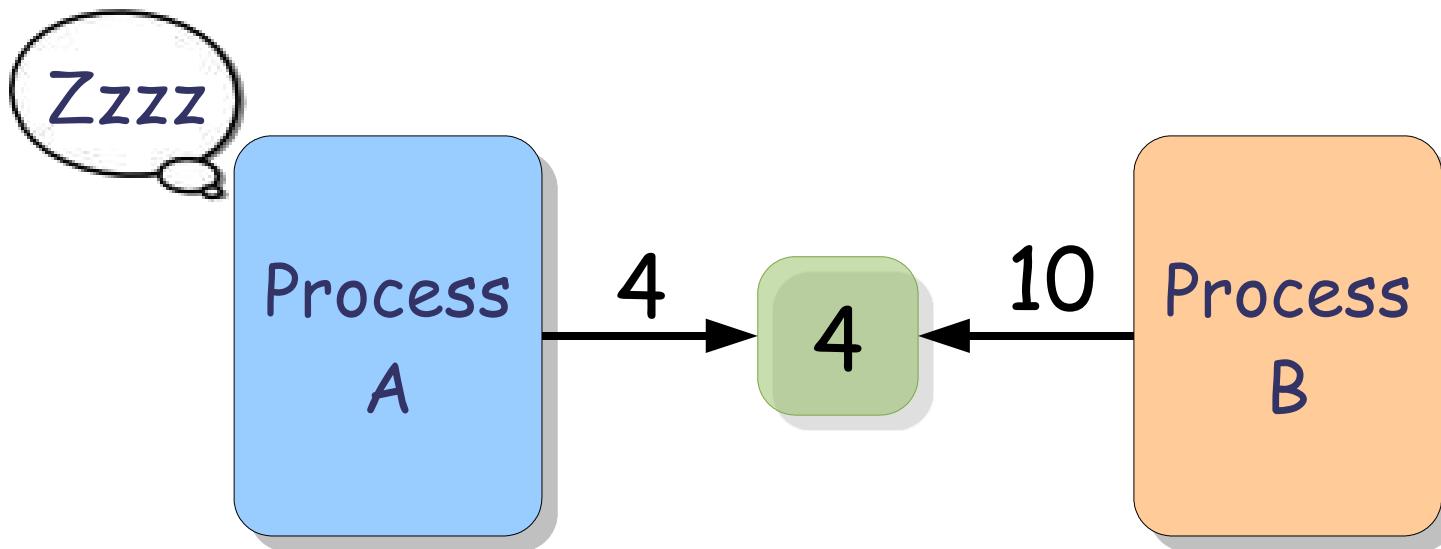
Concurrency is good but ...

- **Atomicity**
 - Non-interference problems (race conditions)
- **Synchronization**
 - Rendez-vous problems (deadlocks, livelocks)
- **Mutual Exclusion**
 - Critical sections problems (starvation)



Atomicity

An operation is said **Atomic** if the result can be observed before its termination

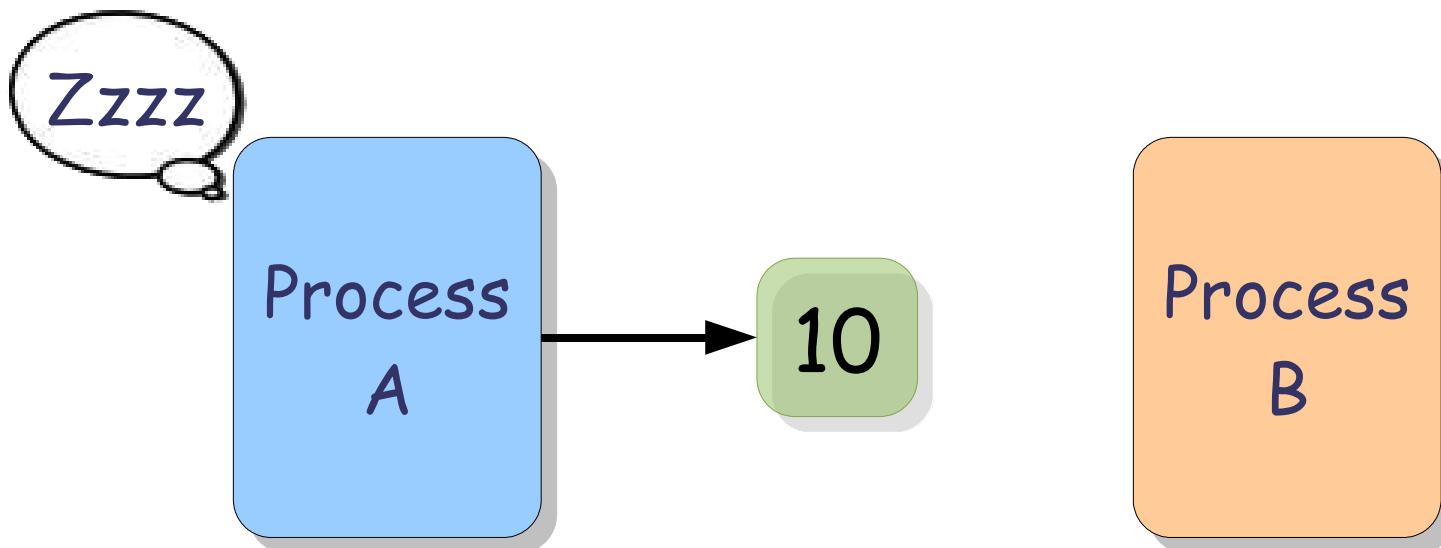


Race Condition



Atomicity

An operation is said **Atomic** if the result can be observed before its termination

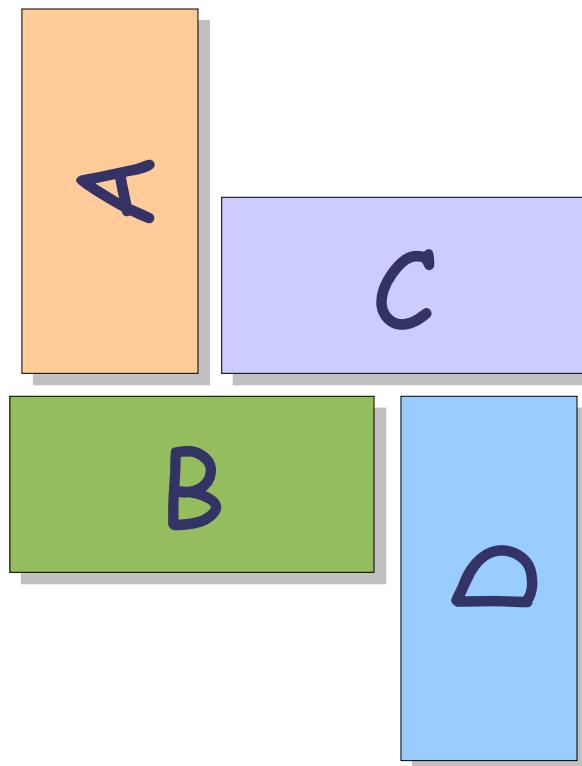


Race Condition

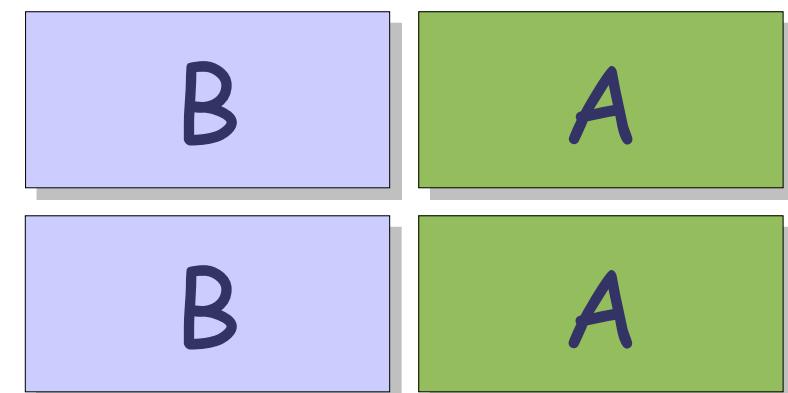


Synchronization

Two processes are **synchronized** when they can exchange some informations



Deadlock

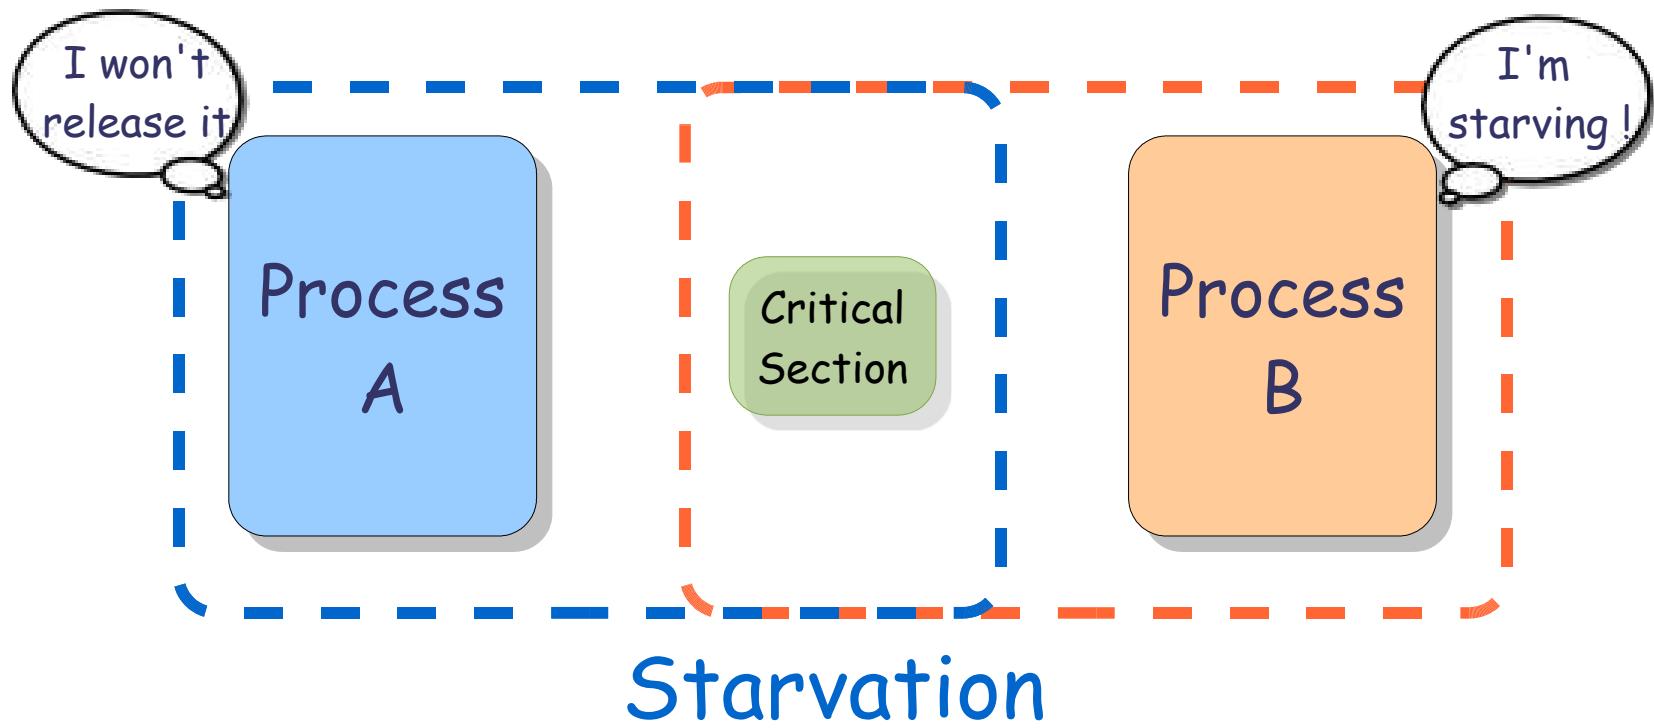


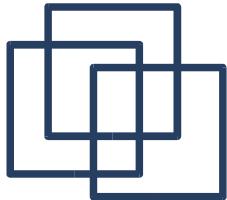
Livelock



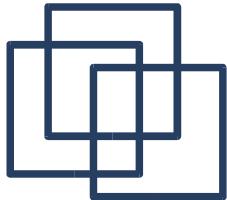
Mutual Exclusion

Two processes are **synchronized** when they can exchange some informations

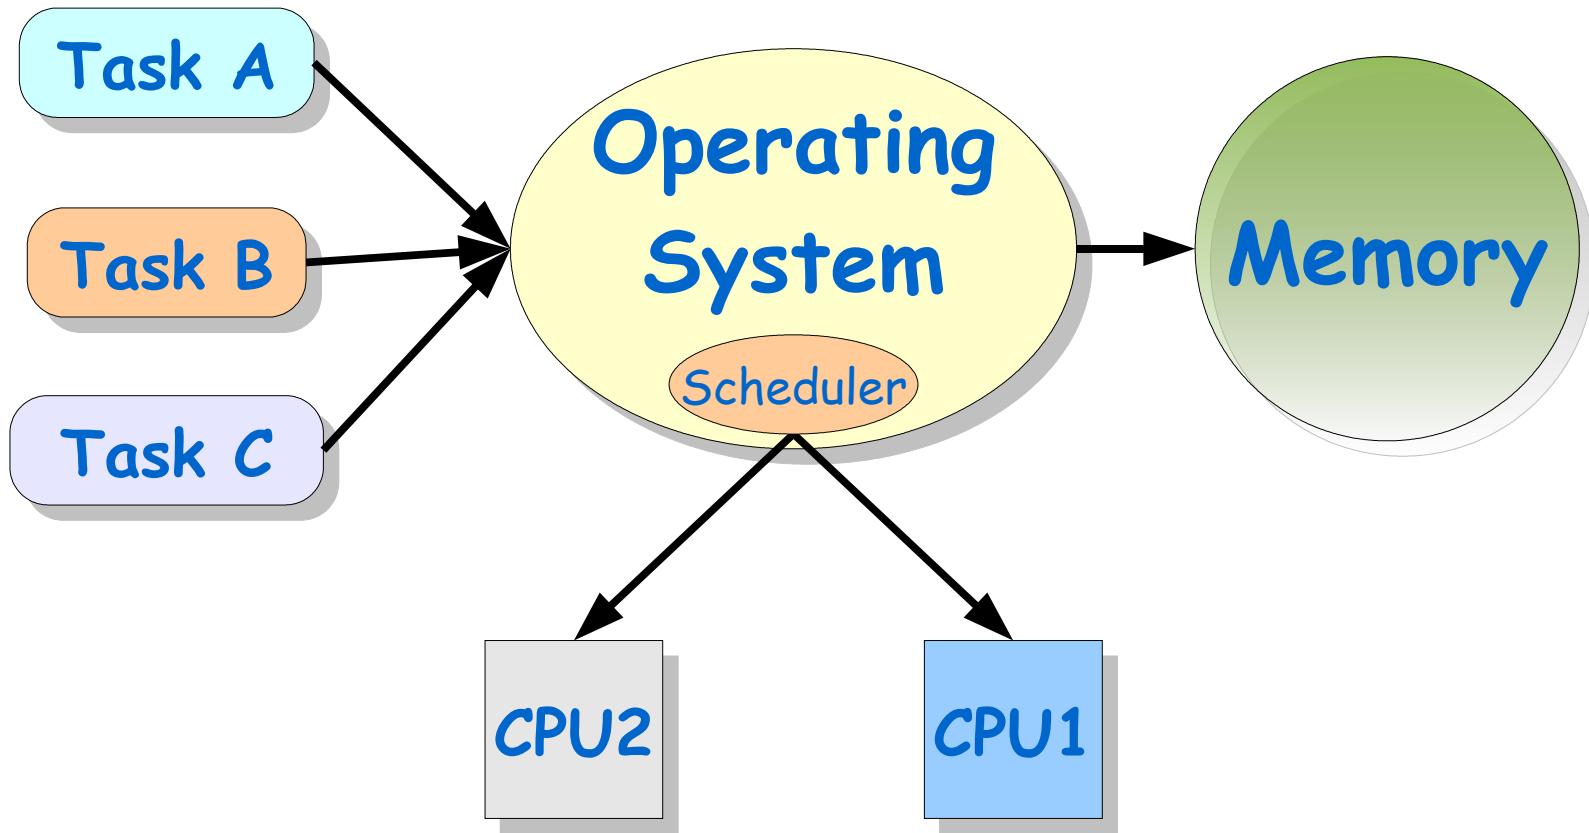




A Process Life



The Role of an OS



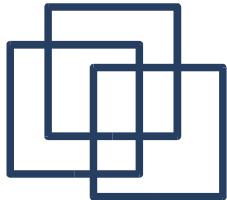
Brings Abstraction from the Hardware !



The Role of the OS

The programmer can assume:

1. Unlimited Resources
(CPU, Memory, ...)
 2. Each task is protected from the others
(Execution, Memory, ...)
 3. Access to the Resources is "fair"
(No starvation induced by the Scheduler)
-



Why Processes ?

- **More Tasks than Processors**

The Scheduler needs to split tasks into smaller units that can be executed on the processor(s) one after one.

- **Making it easy for the Programmer**

The Scheduler make believe to each task that it is the only one on the machine (errors in a task won't interfere with others).

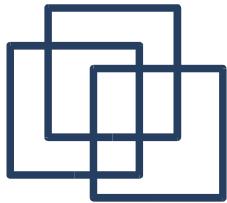
- **Prioritizing Programs**

The Scheduler provides some control from user-space on how often is executed one task and, more generally, what resources does it takes.

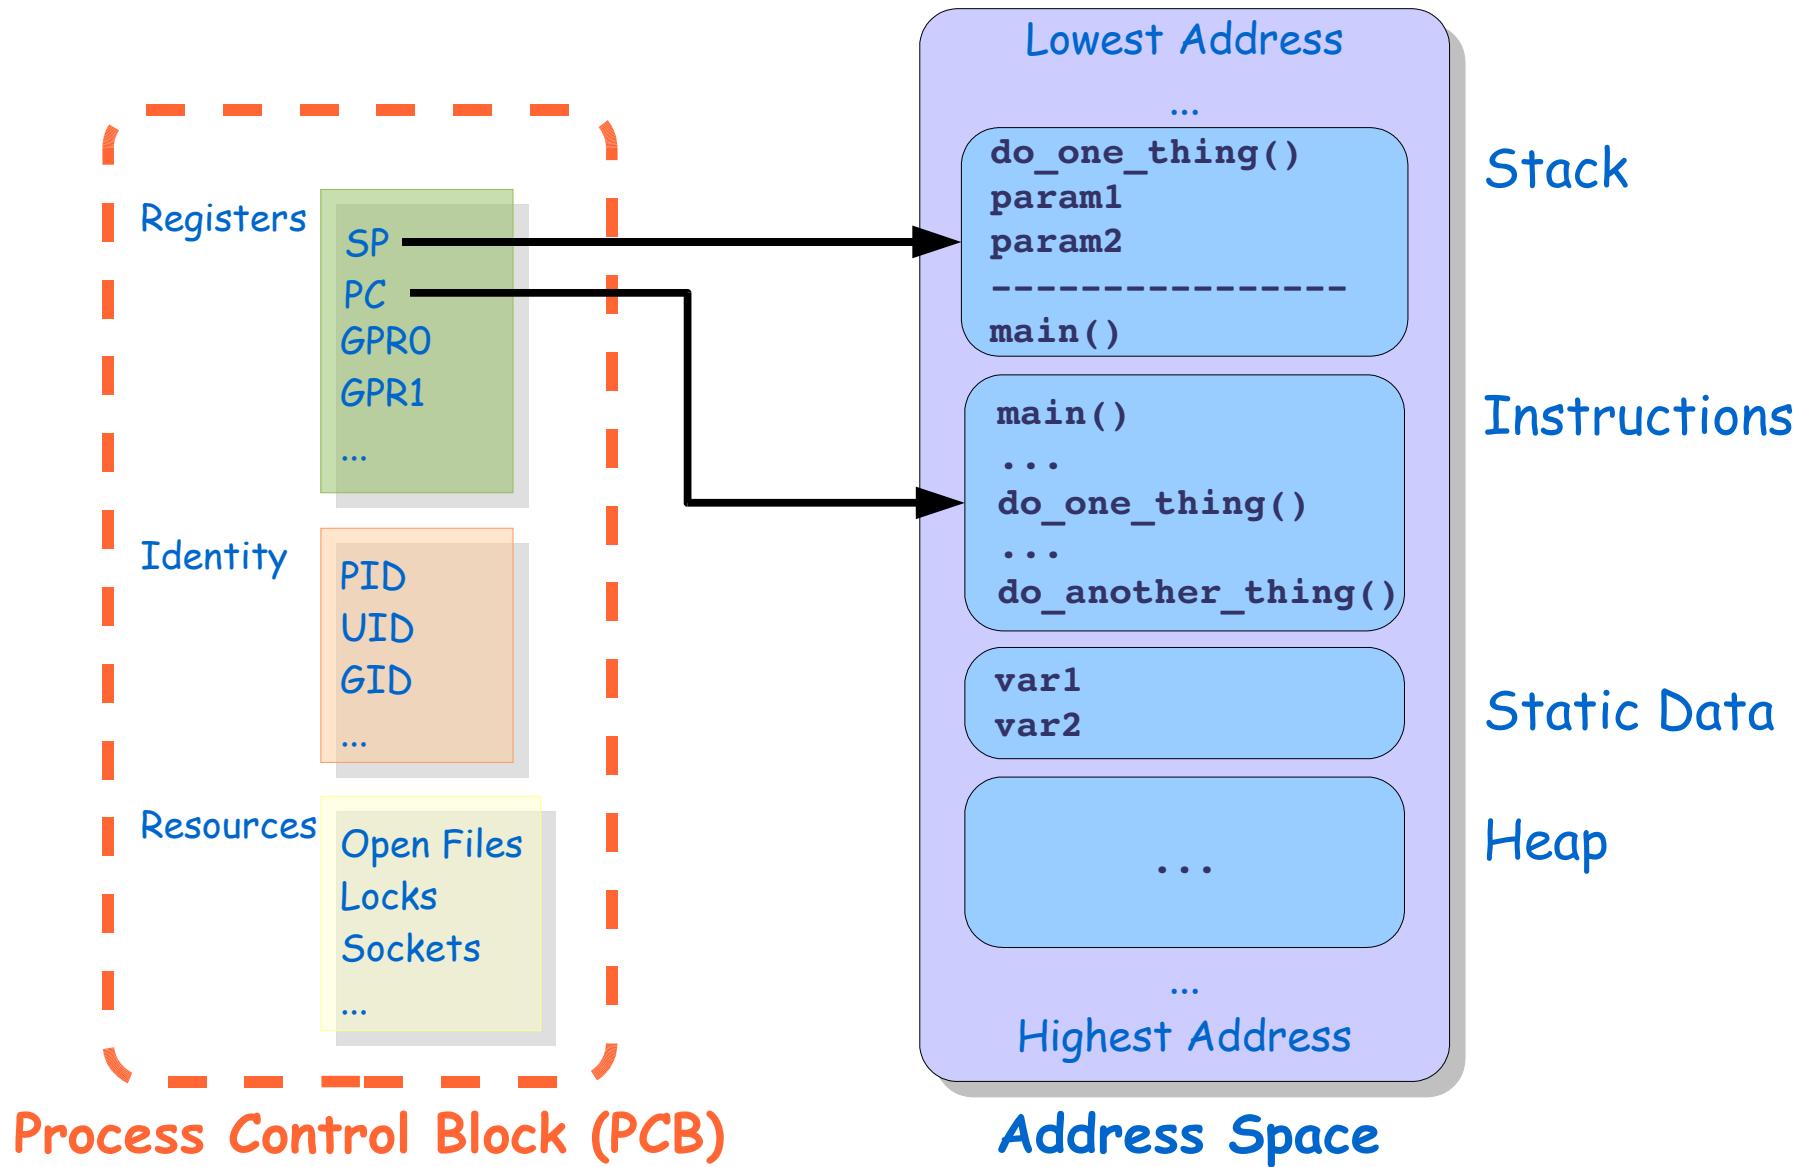


What is a Process ?

- One of the two main abstraction of Unix
(the other one is "everything is a file")
- A Process is the biggest processing unit
that can be scheduled
(the smallest are the threads)
- A Process always spawn from another one
(except the process `init`)



Process Internal





Process Control Block (PCB)

- **PID (Process ID)**: Task's unique process ID, which periodically wraps, though never restarting at zero.
- **PPID (Parent Process ID)**: Process ID of a task's parent.
- **UID (User ID)**: Effective user ID of the task's owner.
- **USER (User Name)**: Effective user name of the task's owner.
- **GROUP (Group Name)**: Effective group name of the task's owner.
- **PR (Priority)**: Priority of the task.
- **NI (Nice value)**: Nice value of the task. A negative nice value means higher priority, whereas a positive nice value means lower priority.



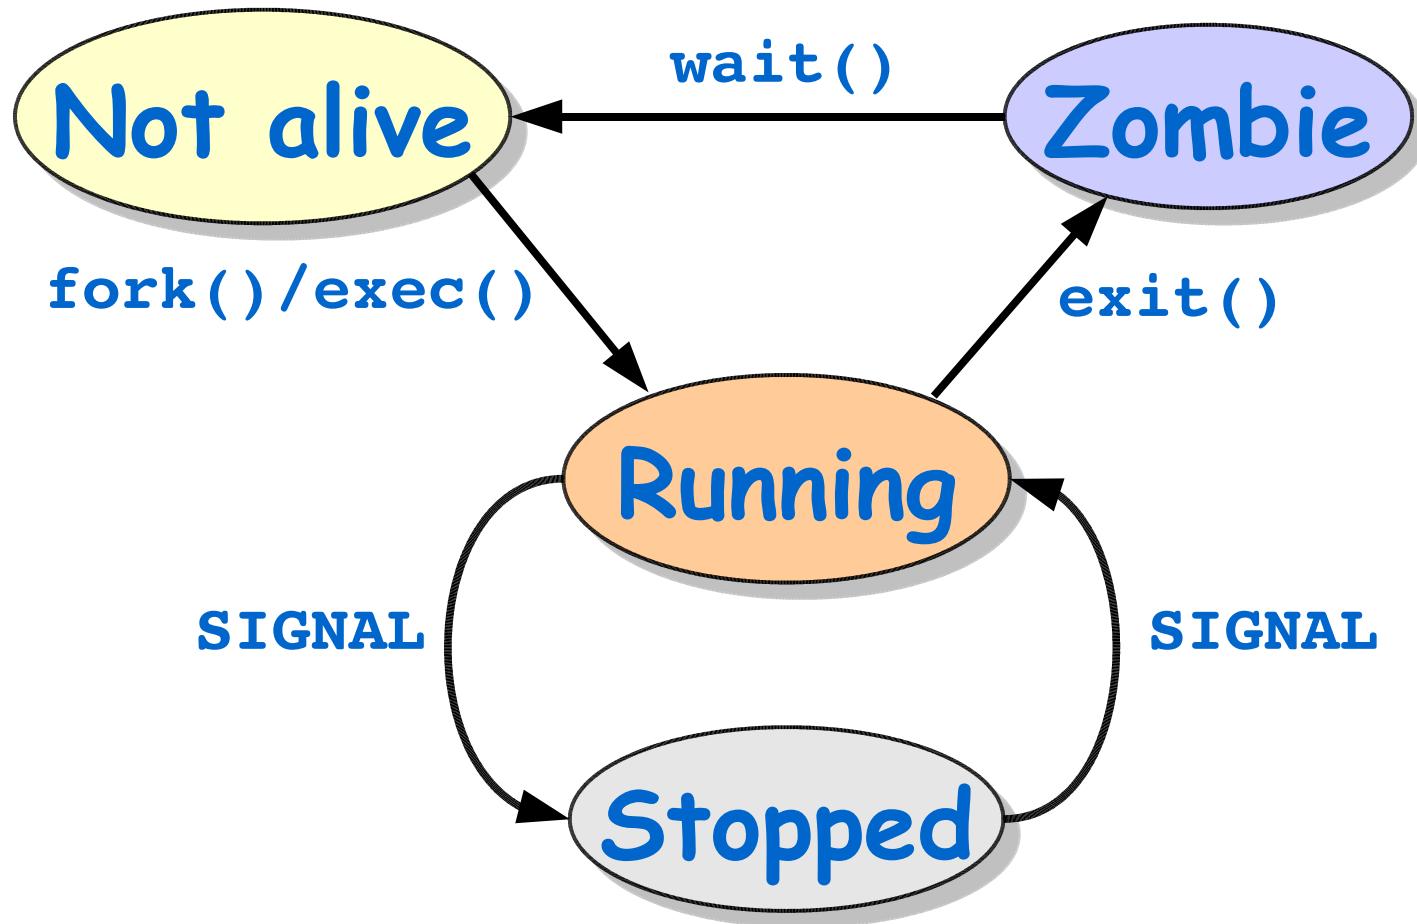
The command “top”

```
top - 00:36:43 up 16:49,  5 users,  load average: 0.91, 0.60, 0.32
Tasks: 76 total,  1 running, 75 sleeping,  0 stopped,  0 zombie
Cpu(s): 6% us, 1.0% sy, 0.0% ni, 92% id, 0.0% wa, 0.0% hi, 0.0% si
Mem: 507576k total, 495948k used,   11628k free, 20016k buffers
Swap: 497972k total,    4184k used, 493788k free, 212796k cached
```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
4299	root	5	-10	202m	69m	6884	S	2.7	14.0	15:55.13	XFree86
5363	fleury	15	0	93896	42m	10m	S	1.7	8.5	8:20.37	rhythmbox
4671	fleury	15	0	30752	14m	7780	S	1.3	2.9	0:56.37	terminal
4665	fleury	16	0	11556	7476	5900	S	0.7	1.5	0:46.25	metacity
1	root	16	0	1500	516	456	S	0.0	0.1	0:00.47	init
2	root	34	19	0	0	0	S	0.0	0.0	0:00.00	ksoftirqd/0
3	root	5	-10	0	0	0	S	0.0	0.0	0:00.73	events/0
4	root	5	-10	0	0	0	S	0.0	0.0	0:00.00	khelper
16	root	15	-10	0	0	0	S	0.0	0.0	0:00.00	kacpid
107	root	5	-10	0	0	0	S	0.0	0.0	0:00.10	kblockd/0
120	root	15	0	0	0	0	S	0.0	0.0	0:00.00	khubd
197	root	15	0	0	0	0	S	0.0	0.0	0:00.02	pdflush
198	root	15	0	0	0	0	S	0.0	0.0	0:00.03	pdflush
200	root	15	-10	0	0	0	S	0.0	0.0	0:00.00	aio/0



Process States





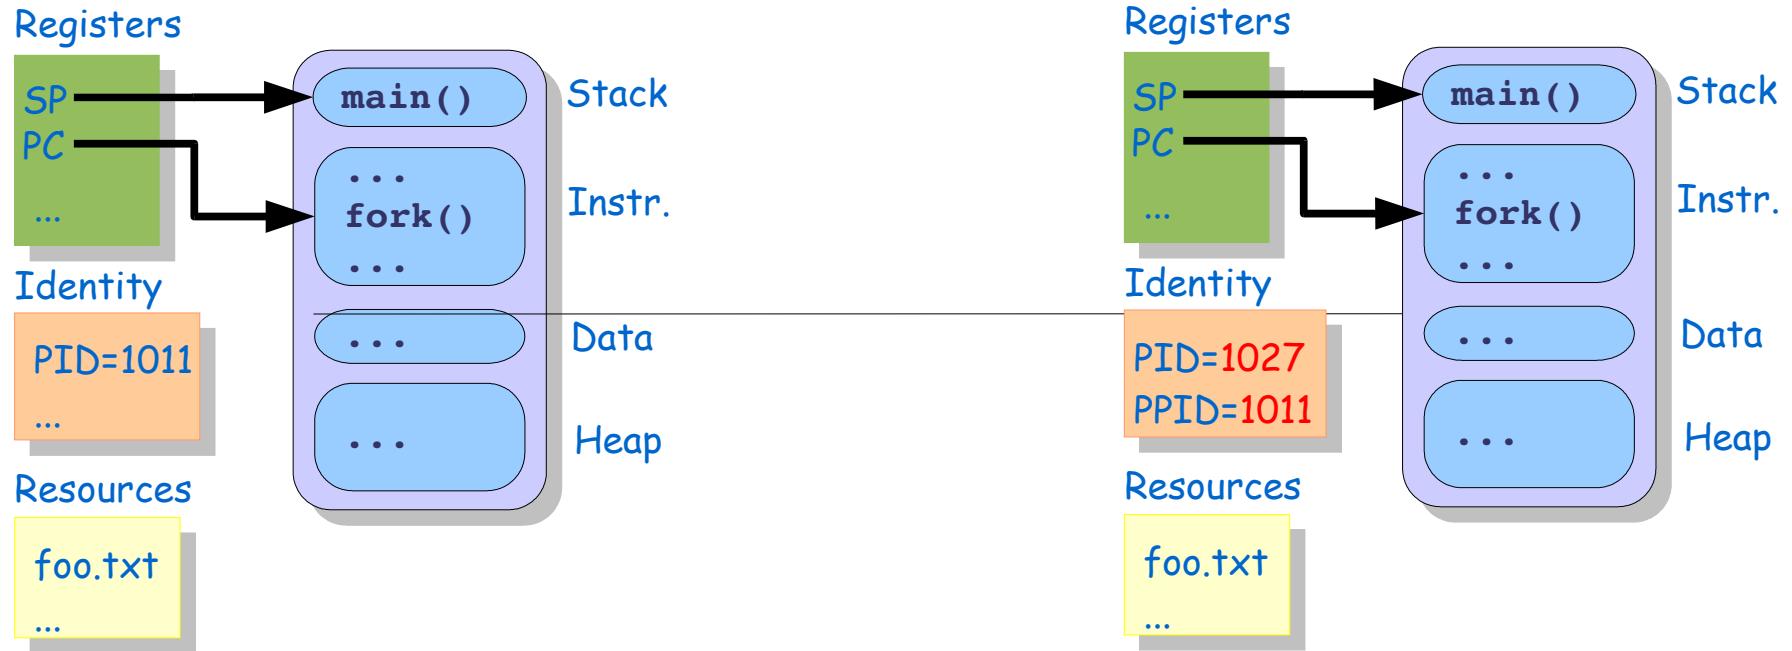
Creation of a Process

- System Call `fork()`:
 - Creates (nearly) identical copy of process
 - Return value different for child/parent

- System Call `exec()`:
 - Over-write with new process memory
 - Return value is 0 for success and 1 for failure



Creation of a Process (fork)



Return Value:

- In Parent Process: “**Child Process ID**” or “**-1**” (on failure)
- In Child Process: “**0**” (always)



fork()

```
#include <stdlib.h>
#include <stdio.h>
#include <unistd.h>
#include <sys/types.h>

int main() {
    pid_t pid;

    switch(pid = fork())
    {
        case -1: /* Failure */
            perror("forking");
            exit(1);

        case 0: /* Child code */
            printf("Child is running\n");
            exit(0);

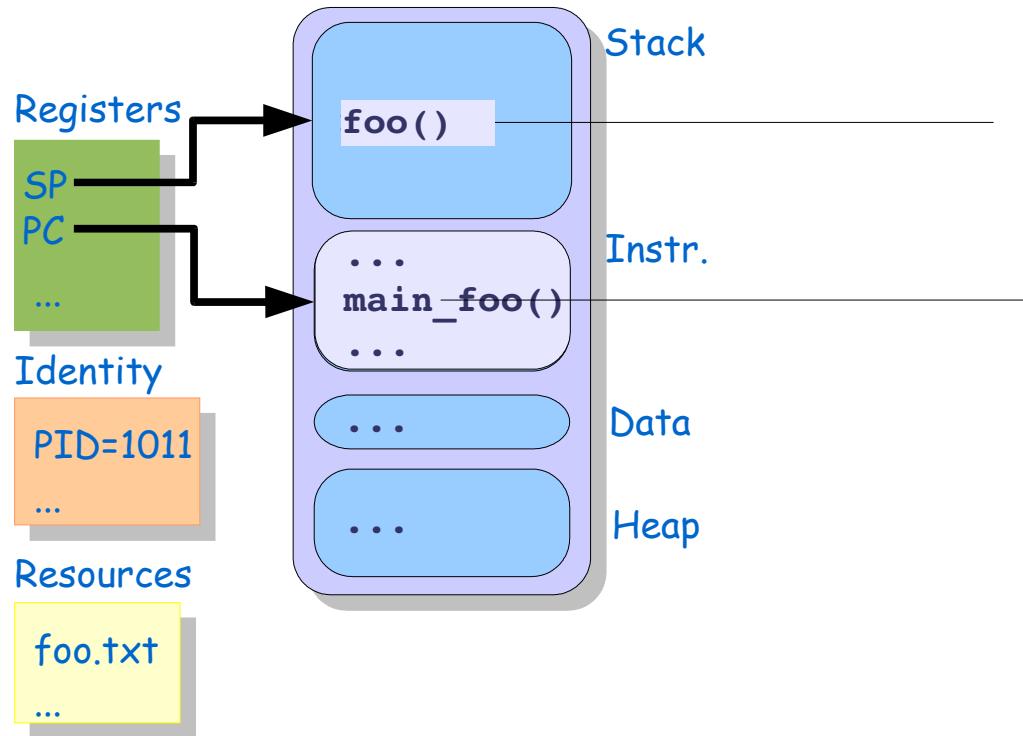
        default: /* Parent Code */
            printf("Parent is running\n");
            exit(0);
    }
}
```

```
[fleury@hermes]$ ./forking
Child is running
Parent is running
[fleury@hermes]$
```

Note: The Linux kernel runs
always the child first.



Creation of a Process (exec)



Return Value:

- "-1" (on failure) and "errno" is set to the error number
- Does not return on success



execve()

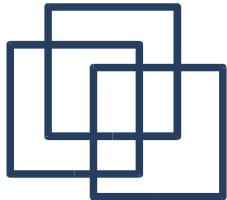
```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
```

```
int main() {
    char *cmd[] = {"ls", "-lh", (char *)NULL};
    char *env[] = {"HOME=/usr/home", "LOGNAME=home", (char *)0};

    if (execve("/bin/ls", cmd, env)) {
        perror("foo");
        exit(1); /* Failure */
    }

    printf("I'm still alive !\n");
    exit(0);
}
```

Note: This code is never executed.

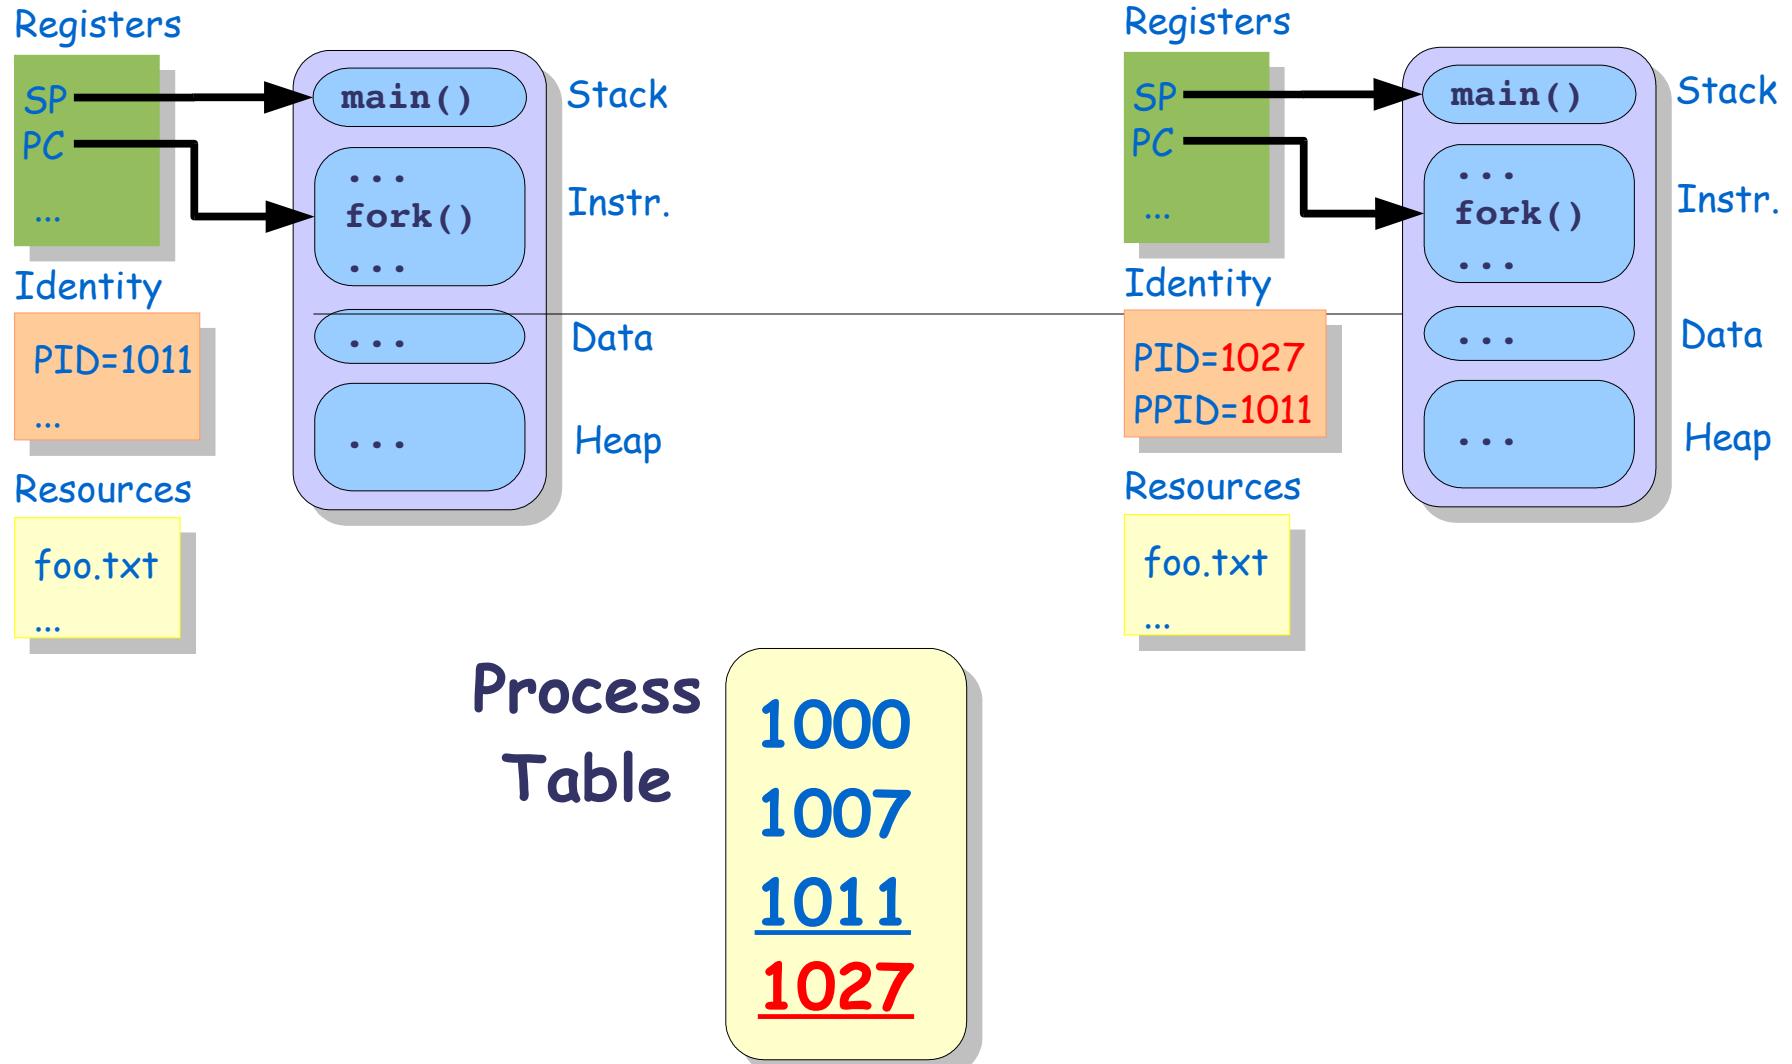


The `exec*`() Familly

- `execve()`: Original system call, all other `exec*`() functions are just front-end to it.
- Other `exec*`() functions are:
`exec1()`, `execlp()`, `execle()`, `execv()`, `execvp()`.
- What's the difference ?
 - l/v = Arguments are given as a "list" or a "vector"
 - p/e = Environment is used as such (p) or given as an argument (e).
 - Examples:
 - `execlp("li", "li", "-al", 0);`
 - `execl("/usr/bin/sh", "sh", "-c", "li -l *.c", 0);`

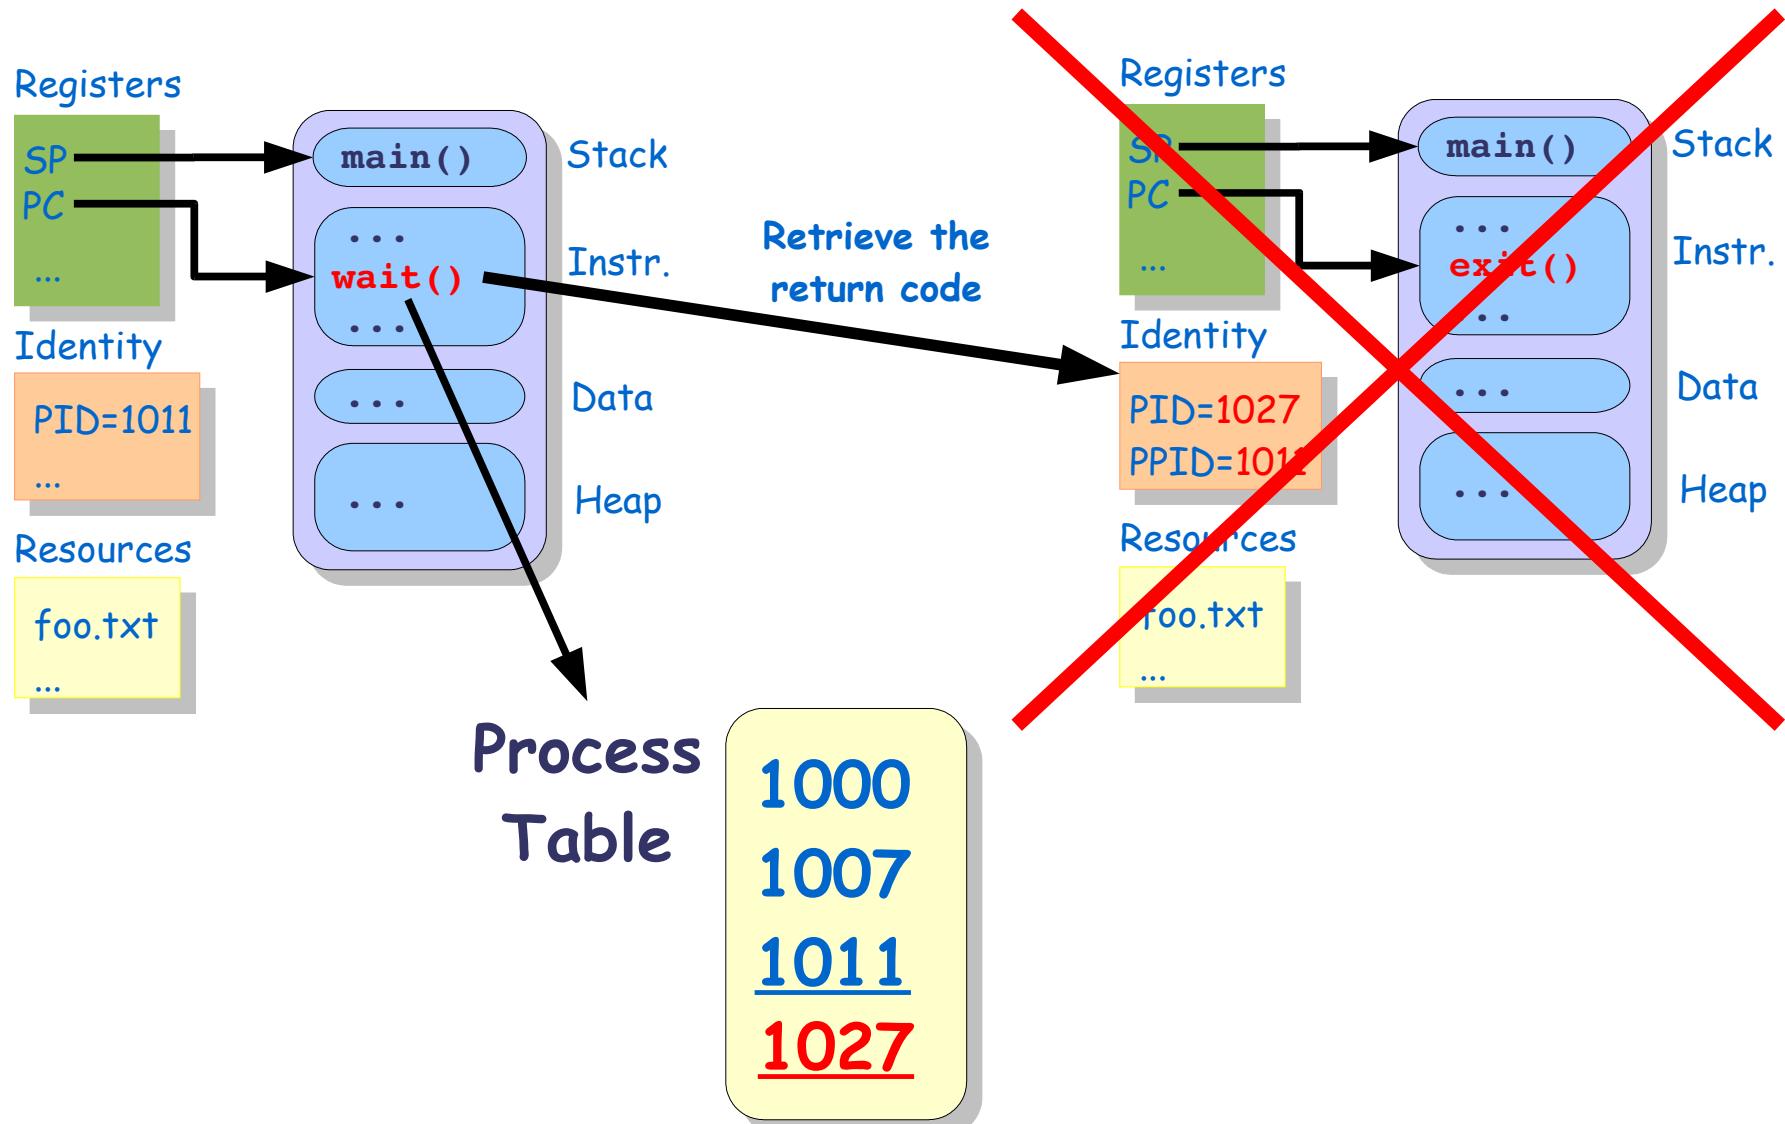


Termination of a Process (exit)





Termination of a Process (exit)





wait()

```
#include <stdlib.h>
#include <stdio.h>
#include <unistd.h>
#include <sys/types.h>

int main() {
    pid_t pid;

    switch(pid = fork())
    {
        case -1: /* Failure */
            perror("waiting");
            exit(1);

        case 0: /* Child code */
            printf("Child is running\n");
            exit(0);

        default: /* Parent Code */
            printf("Parent is running\n");
            while(1);
            exit(0);
    }
}
```

```
[fleury@hermes]$ ./waiting
Child is running
Parent is running

[3]+  Stopped                  ./waiting
[fleury@hermes]$ ps a | grep waiting
PID  TTY      STAT   TIME CMD
24859 pts/3    T      0:01 ./waiting
24860 pts/3    Z      0:00 [waiting] <defunct>
```

- D Uninterruptible sleep (usually IO)
- R Running or runnable (on run queue)
- S Interruptible sleep
(waiting for an event to complete)
- T Stopped, either by a job control signal
or because it is being traced.
- X dead (should never be seen)
- Z Defunct ("zombie") process,
terminated but not reaped by its parent.



wait()

```
#include <stdlib.h>
#include <stdio.h>
#include <unistd.h>
#include <sys/types.h>
#include <sys/wait.h>

int main() {
    pid_t pid;

    switch(pid = fork())
    {
        case -1: /* Failure */
            perror("waiting");
            exit(1);
        case 0: /* Child code */
            printf("Child is running\n");
            exit(0);
        default: /* Parent Code */
            printf("Parent is running\n");
            while (pid != wait(&status));
            printf("The Child %i has returned the value %i\n", pid, status/256);
            exit(0);
    }
}
```

```
[fleury@hermes]$ ./waiting
Child is running
Parent is running
The Child 25543 has returned the value 0
[fleury@hermes]$
```

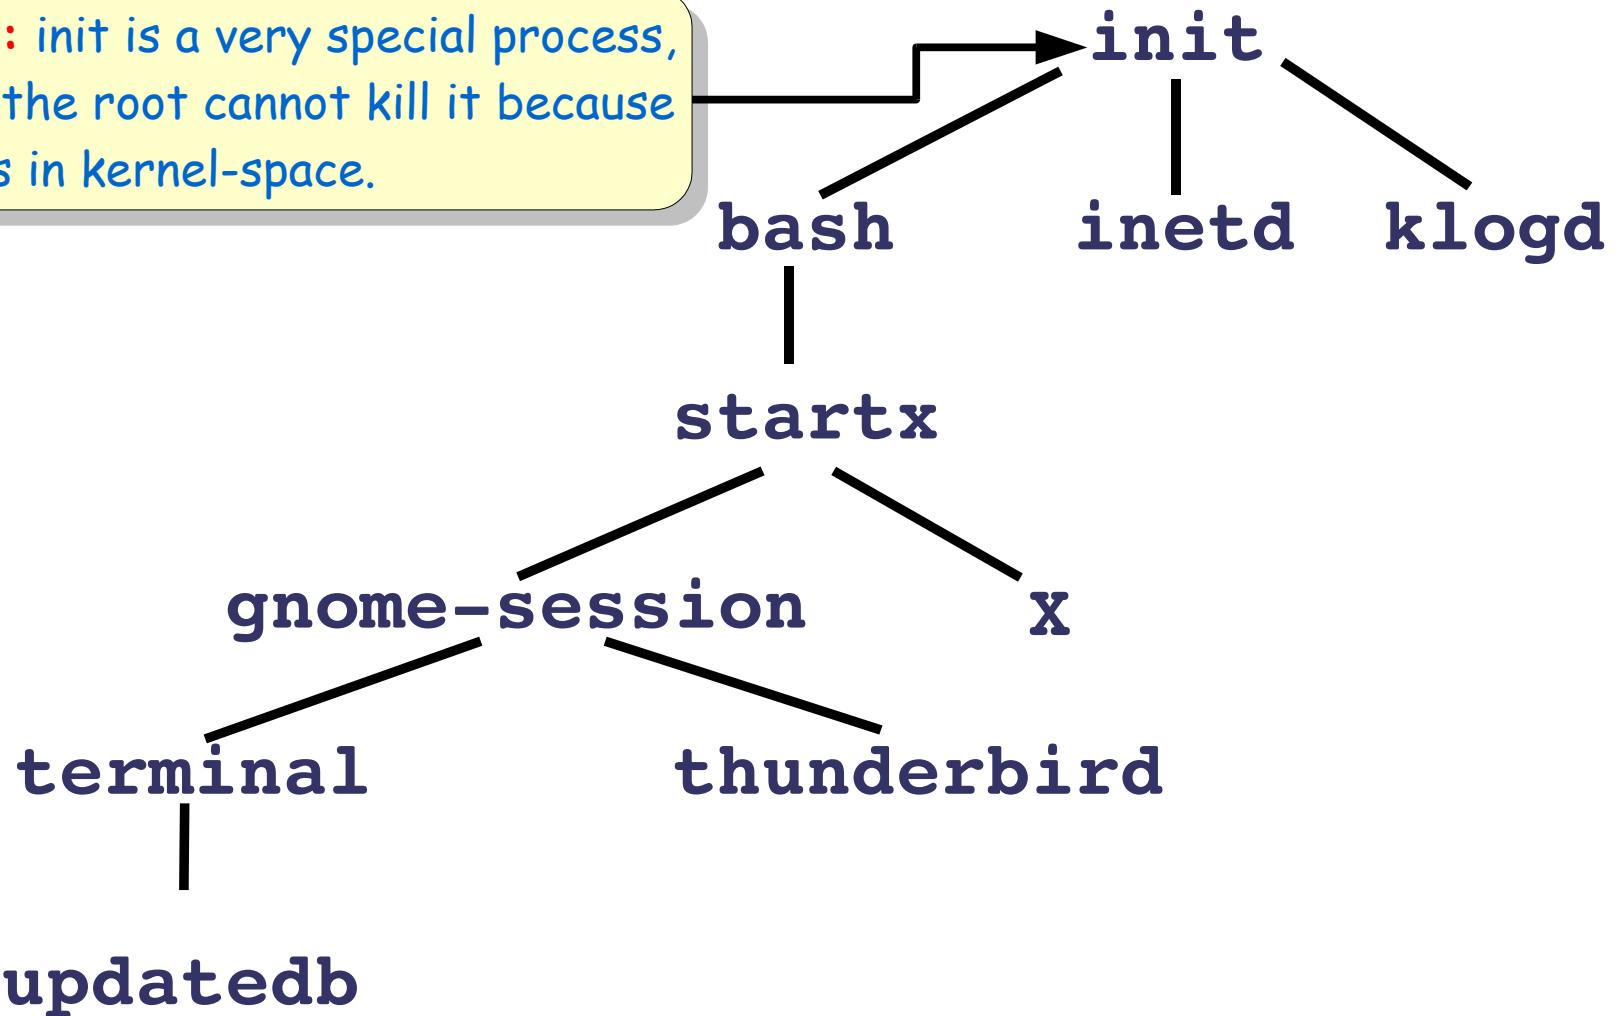
Note: We are waiting for a precise child (pid) but we have only one, this could be avoided.

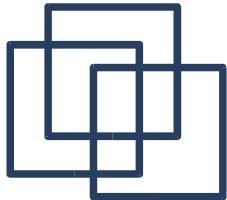




Process Tree

Note: init is a very special process, even the root cannot kill it because it lies in kernel-space.





Interlude: Processes in (nut)Shell



Basic Commands

- **bg (^Z or &):**
Send the current job in the background
 - **fg:**
Run the background job in the foreground
 - **jobs:**
List all the jobs present on the shell
 - **kill (^C):**
Terminate a job
 - **wait:**
Wait for the termination of a job
-



Demonstration of bg, fg, jobs, kill

- **bg:** Put several jobs in the background
- **jobs:** List them all
- **fg:** Select one and run it in foreground
- **kill:** Send termination signals to some of the background jobs



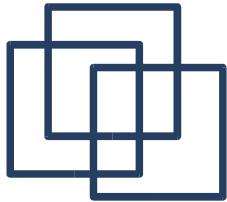
wait (Shell)

```
#!/bin/sh

# Job 1
ls -a &
# PID of Job 1
p1=$!
# Job 2
ls -al &
# Display status of Job 1
wait $p1
echo Job 1 exited with status $?
# Display status of Job 2
wait $!
echo Job 2 exited with status $?
```

```
[fleury@hermes]$ ./test.sh
. .. test.sh
Job 1 exited with status 0
total 16
drwxr-xr-x  2 fleury fleury 4096 Mar 25 11:06 .
drwxr-xr-x  74 fleury fleury 4096 Mar 25 11:05 ..
-rwxr-xr-x  1 fleury fleury 152 Mar 25 11:06 test.sh
Job 2 exited with status 0
[fleury@hermes]$ ./test.sh
total 16
drwxr-xr-x  2 fleury fleury 4096 Mar 25 11:06 .
drwxr-xr-x  74 fleury fleury 4096 Mar 25 11:05 ..
-rwxr-xr-x  1 fleury fleury 152 Mar 25 11:06 test.sh
. .. test.sh
Job 1 exited with status 0
Job 2 exited with status 0
[fleury@hermes]$
```

Note: The order of execution might change.



Process Management



The `get*`() Family

- `getpid()`: Get process ID
- `getppid()`: Get parent process ID
- `getuid()`: Get user ID
- `geteuid()`: Get effective user ID (ignoring set ID calls)
- `getgid()`: Get group ID
- `getegid()`: Get effective group ID (ignoring set ID calls)
- `getresuid()`: Get real, effective and saved user ID
- `getresgid()`: Get real, effective and saved group ID
- `getgroups()`: Get the list of groups to which belong the user



The command “id”

```
[fleury@hermes]$ id  
uid=1000(fleury) gid=1000(fleury) groups=29(audio),1000(fleury)  
[fleury@hermes]$ cp /bin/sh .  
[fleury@hermes]$ chmod +s sh  
[fleury@hermes]$ su -  
Password:  
[root@hermes]$ id  
uid=0(root) gid=0(root) groups=0(root)  
[root@hermes]$ ls -l ~fleury/sh  
-rwsr-sr-x 1 fleury fleury 667180 Mar 26 17:26 /home/fleury/sh  
[root@hermes]$ ~fleury/sh  
[root@hermes]$ id  
uid=0(root) gid=0(root) euid=1000(fleury) egid=1000(fleury) groups=0(root)  
[root@hermes]$ exit  
[root@hermes]$ id  
uid=0(root) gid=0(root) groups=0(root)  
[root@hermes]$ exit  
[fleury@hermes]$
```



getgroups()

```
#include <stdlib.h>
#include <stdio.h>
#include <unistd.h>
#include <sys/types.h>

int main() {
    int size, i;
    gid_t *gid_table;

    /* Get the size of the table */
    if ((size = getgroups(0, NULL)) < 0) {
        perror("get_groups");
        exit(1);
    }
    /* Memory allocation of the table */
    if ((gid_table = calloc(size, sizeof(gid_t))) == NULL){
        perror("get_groups");
        exit(1);
    }
    /* Get the group list */
    if (!getgroups(size, gid_table)) {
        perror("get_groups");
        exit(1);
    }
    /* Display the list */
    for (i=0; i<size; i++)
        printf("group[%i] = %u\n", i, gid_table[i]);
    free(gid_table);

    exit(0);
}
```

```
[fleury@hermes]$ id
uid=1000(fleury) gid=1000(fleury) groups=29(audio),1000(fleury)
[fleury@hermes]$ ./get_groups
group[0] = 29
group[1] = 1000
[fleury@hermes]$
```



The `set*`() Family

- **`setuid()`/`setgid()`:**

Sets the effective user/group ID of the current process

- **`setresuid()`/`setresgid()`:**

Sets the real user ID, the effective user ID, and the saved (effective) user ID of the current process.

- **`seteuid()`/`setegid()`:**

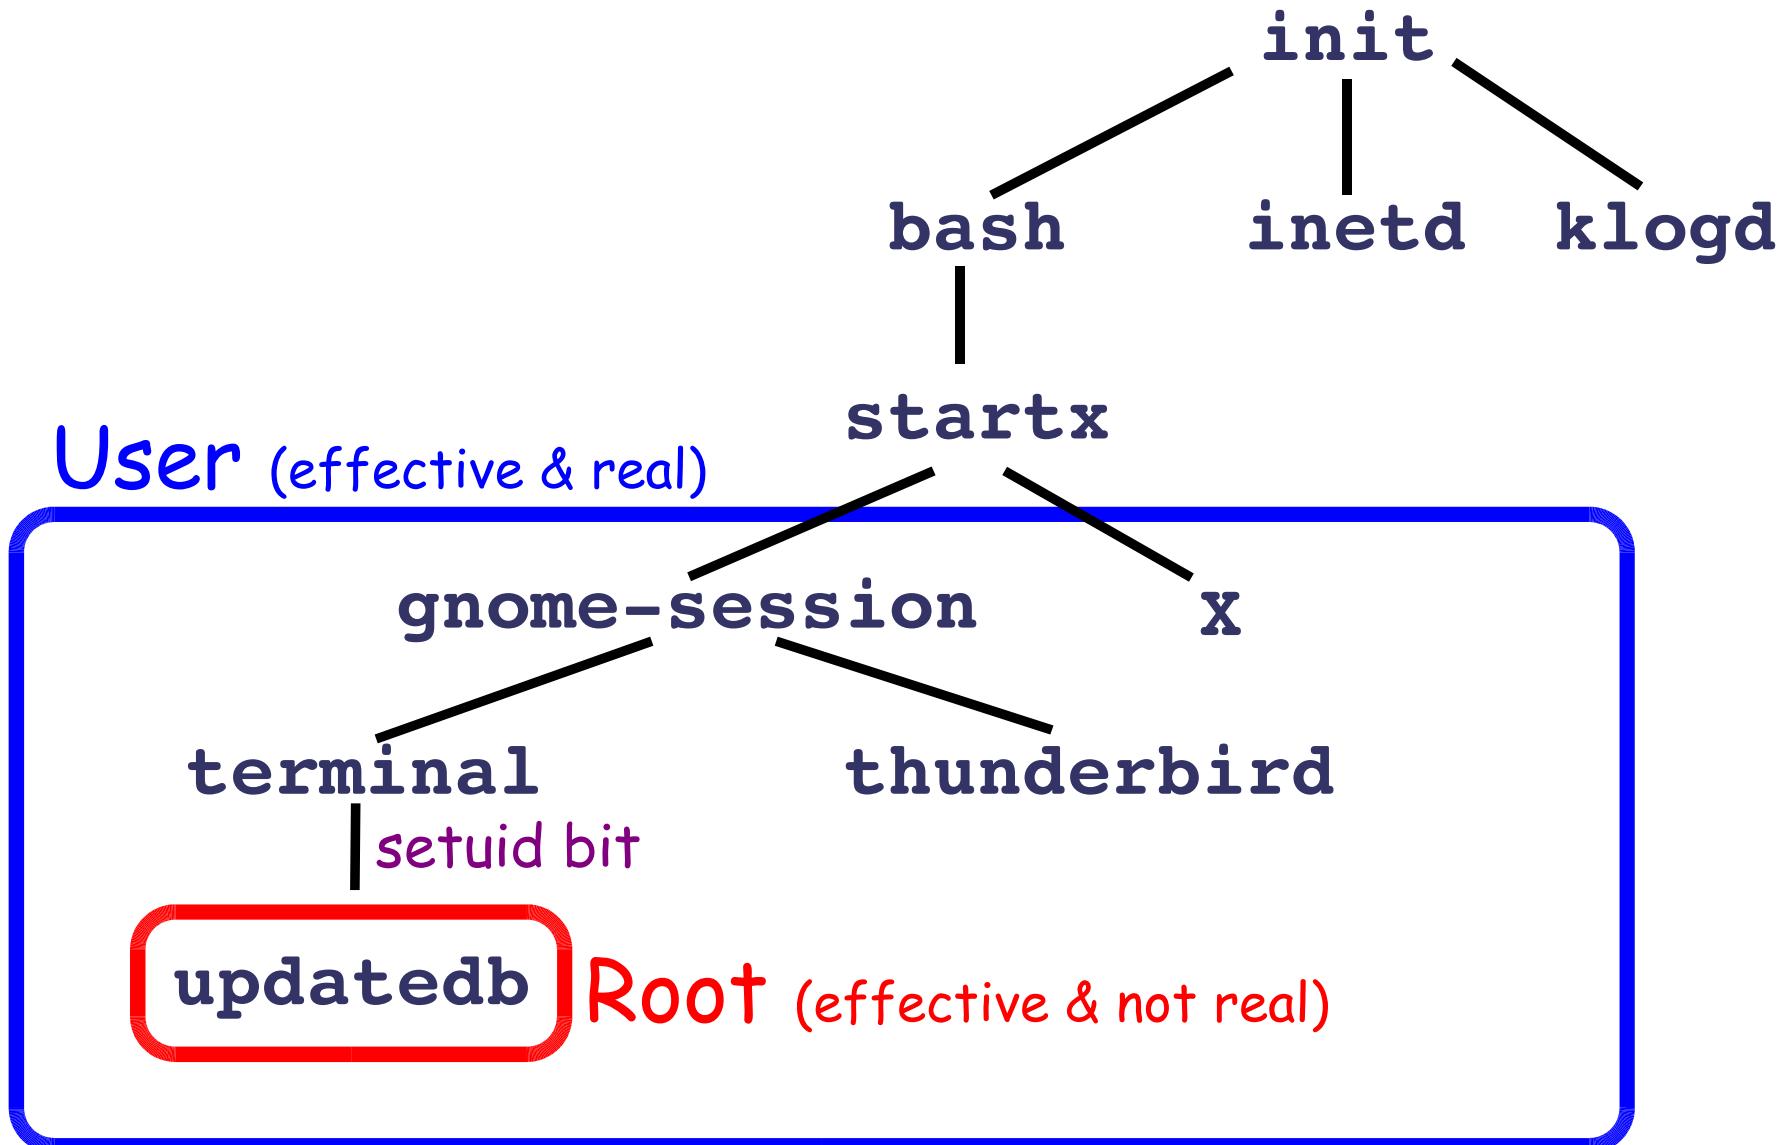
Sets the effective user/group ID of the current process

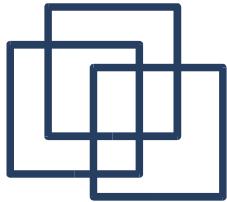
- **`setreuid()`/`setregid()`:**

Sets real and effective user IDs of the current process.



Process Tree



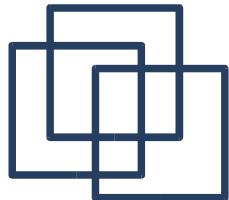


sleep()

- **sleep():**
Sleep for the specified number of seconds
- **usleep():**
Suspend execution for microsecond intervals
- **nanosleep():**
Pause execution for a specified time



Process Scheduling



Getting & Setting Priority

- **nice()**:

Change the nice value of a process. Return the new priority value or "-1" in case of failure

- **renice()**:

Alter priority of running processes

- **getpriority()**:

Get program scheduling priority

- **setpriority()**:

Set program scheduling priority



nice()

```
#include <sys/resource.h>
#include <stdlib.h>
#include <stdio.h>
#include <unistd.h>

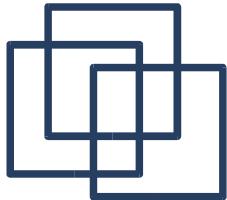
int main() {

    printf("process priority is %i\n",
           getpriority(PRIO_PROCESS, 0));
    printf("process group priority is %i\n",
           getpriority(PRIO_PGRP, 0));
    printf("user priority is %i\n",
           getpriority(PRIO_USER, 0));
    printf("=====\n");
    printf("new nice value: %i\n", nice(3));
    printf("new nice value: %i\n", nice(3));
    printf("new nice value: %i\n", nice(-9));
    printf("=====\n");
    printf("process priority is %i\n",
           getpriority(PRIO_PROCESS, 0));
    printf("process group priority is %i\n",
           getpriority(PRIO_PGRP, 0));
    printf("user priority is %i\n",
           getpriority(PRIO_USER, 0));

    exit(0);
}
```

Note: Priority is a value between -20 and 19.
Only root can go under 0.

```
[fleury@hermes]$ ./priority
process priority is 0
process group priority is 0
user priority is 0
=====
new nice value: 3
new nice value: 6
new nice value: -1
=====
process priority is 6
process group priority is 6
user priority is 0
[fleury@hermes]$ su -c ./priority
Password:
process priority is 0
process group priority is 0
user priority is -10
=====
new nice value: 3
new nice value: 6
new nice value: -3
=====
process priority is -3
process group priority is -3
user priority is -10
```

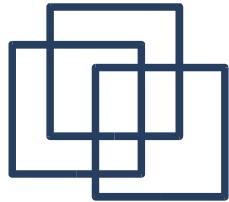


See also ...

- `mlock()` / `munlock()`
 - `mlockall()` / `munlockall()`
 - `sched_get_priority_max()` / `sched_get_priority_min()`
 - `sched_getaffinity()` / `sched_setaffinity()`
 - `sched_getparam()` / `sched_setparam()`
 - `sched_getscheduler()` / `sched_setscheduler()`
 - `sched_rr_get_interval()`
 - `sched_yield()`
 - `capabilities()`
 - ...
-



Questions ?



Next Week

- Signals in Unix
- Inter-process communication
 - files
 - pipes
 - named pipes
 - sharing memory chunks
 - FIFO
 - Semaphores